Generative AI Interview Questions:

Sujatha Mudadla · Follow
10 min read · May 20, 2024

1. **What is the difference between a Generative Adversarial Network (GAN) and a Variational Autoencoder (VAE)?**

Answer: GANs consist of a generator and a discriminator, whereas VAEs consist of an encoder and a decoder. GANs are more suitable for generating new samples, whereas VAEs are more suitable for dimensionality reduction and feature learning.

2. How do you train a GAN?

Answer: A GAN can be trained using an alternating optimization approach, where the generator and discriminator are trained alternately.

3. What is the concept of mode collapse in GANs?

Answer: Mode collapse occurs when a GAN generates limited variations of the same output, rather than exploring the full range of possibilities.

4. How do you evaluate the performance of a GAN?

Answer: GAN performance can be evaluated using metrics such as **Frechet Inception Distance (FID), Inception Score (IS), and Visual Turing Test (VTT).**

5. What is the difference between a conditional GAN and an unconditional GAN?

Answer: A conditional GAN generates samples based on a specific condition or class, whereas an unconditional GAN generates samples without any specific condition or class.

6. How do you implement a text-to-image synthesis model?

Answer: Text-to-image synthesis can be implemented using a GAN, where the generator takes a text input and generates an image, and the discriminator evaluates the generated image.

7. What is the concept of disentanglement in generative models?

Answer: Disentanglement refers to the ability of a generative model to separate the underlying factors of variation in the data, such as object pose, shape, and color.

Answer: Style transfer can be performed using a GAN, where the generator takes a content image and a style image as input, and generates an output image that combines the content and style.

9. What is the concept of attention in generative models?

Answer: Attention refers to the ability of a generative model **to focus on specific parts of the input** data when generating output.

10. How do you implement a video generation model using generative AI?

Answer: Video generation **can be implemented using a GAN**, where the generator takes a sequence of images as input and generates a video, and the discriminator evaluates the generated video.

11. What is the difference between a generative model and a discriminative model?

Answer: A generative model models the underlying distribution of the data, whereas a **discriminative model** models the probability of a specific output given the input data. **Generative models are used for tasks such as data augmentation, whereas discriminative models are used for tasks such as classification.**

12. How do you handle class imbalance in a classification problem?

Answer: There are several ways to handle class imbalance, including oversampling the minority class, undersampling the majority class, using class weights, and using metrics such as F1-score and AUC-ROC instead of accuracy.

13. What is the difference between a correlation and a causation?

Answer: Correlation refers to the statistical relationship between two variables, whereas causation refers to the cause-and-effect relationship between two variables. Correlation does not imply causation.

14. How do you evaluate the performance of a regression model?

Answer: Regression model performance can be evaluated using metrics such as Mean Squared Error (MSE), Mean Absolute Error (MAE), R-Squared, and Coefficient of Determination.

15. What is the purpose of feature scaling in machine learning?

Answer: Feature scaling is used to normalize the features to a common scale, which can improve the performance of machine learning algorithms by reducing the effect of features with large ranges.

16. How do you handle missing values in a dataset?

Answer: There are several ways to handle missing values, including mean/median imputation, regression imputation, and listwise deletion.

17. What is the difference between a supervised and an unsupervised learning algorithm?

Answer: Supervised learning algorithms are trained on labeled data to make predictions, whereas unsupervised learning algorithms are trained on unlabeled data to discover patterns or relationships.

18. How do you implement a recommendation system?

Answer: A recommendation system can be implemented using collaborative filtering, content-based filtering, or hybrid approaches.

19. What is the concept of overfitting in machine learning?

Answer: Overfitting occurs when a model is too complex and performs well on the training data but poorly on new, unseen data.

20. How do you perform hyperparameter tuning in machine learning?

Answer: Hyperparameter tuning can be performed using grid search, random search, Bayesian optimization, or gradient-based optimization.

21. What is the difference between a decision tree and a random forest?

Answer: A decision tree is a single tree-based model, whereas a random forest is an ensemble of multiple decision trees.

22. How do you handle high-dimensional data in machine learning?

Answer: High-dimensional data can be handled using techniques such as dimensionality reduction (e.g. PCA, t-SNE), feature selection, and regularization.

23. What is the concept of clustering in machine learning?

Answer: Clustering is a type of unsupervised learning where similar data points are grouped together based on their features.

24. How do you implement a natural language processing (NLP) model?

Answer: NLP models can be implemented using techniques such as tokenization, word embeddings, and recurrent neural networks (RNNs).

25. What is the difference between a convolutional neural network (CNN) and a recurrent neural network (RNN)?

Answer: CNNs are used for image and signal processing tasks, whereas RNNs are used for sequential data tasks such as language modeling and machine translation.

26. How do you evaluate the performance of a clustering model?

Answer: Clustering model performance can be evaluated using metrics such as silhouette score, calinski-harabasz index, and davies-bouldin index.

27. What is the concept of transfer learning in machine learning?

Answer: Transfer learning is the process of using a pre-trained model as a starting point for a new task, which can improve performance and reduce training time.

28. How do you implement a reinforcement learning model?

Answer: Reinforcement learning models can be implemented using techniques such as **Q-learning, SARSA, and deep reinforcement learning.**

29. What is the difference between a supervised and a self-supervised learning algorithm?

Answer: Supervised learning algorithms are trained on labeled data, whereas self-supervised learning algorithms are trained on unlabeled data using techniques such as **autoencoders and contrastive learning.**

30. How do you handle outliers in a dataset?

Answer: Outliers can be handled using techniques such as data transformation, outlier detection, and robust regression.

31. What is the difference between a GAN and a VAE in terms of training objectives?

Answer: GANs are trained using an adversarial objective, whereas VAEs are trained using a reconstruction objective.

32. How do you handle the problem of non-differentiable objectives in GANs?

Answer: Non-differentiable objectives can be handled using techniques such as the **REINFORCE algorithm and the Gumbel-softmax trick.**

33. What is the concept of batch normalization in GANs?

Answer: Batch normalization is a technique used to normalize the input data for each layer in a GAN, which can improve stability and performance.

34. How do you implement a GAN for image-to-image translation?

Answer: Image-to-image translation can be implemented using a GAN, where the generator takes an input image and generates an output image, and the discriminator evaluates the generated image.

35. What is the concept of cycle consistency in GANs?

Answer: Cycle consistency refers to the idea that a GAN should be able to translate an input image back to its original form, which can improve performance and stability.

36. How do you evaluate the performance of a VAE?

Answer: VAE performance can be evaluated using metrics such as reconstruction loss, KL divergence, and mutual information.

37. What is the difference between a VAE and a normalizing flow?

Answer: VAEs are trained using a probabilistic objective, whereas normalizing flows are trained using a deterministic objective.

38. How do you implement a normalizing flow?

Answer: Normalizing flows can be implemented using techniques such as planar flows and radial flows.

39. What is the concept of invertible neural networks?

Answer: Invertible neural networks are neural networks that can be inverted, which can be used for tasks such as density estimation and sampling.

40. How do you implement a GAN for data augmentation?

Answer: Data augmentation can be implemented using a GAN, where the generator takes an input image and generates a new image with augmented features, and the discriminator evaluates the generated image.

41. What is the concept of conditional independence in GANs?

Answer: Conditional independence refers to the idea that the generator and discriminator in a GAN should be conditionally independent given the input data.

42. How do you handle the problem of mode collapse in VAEs?

Answer: Mode collapse can be handled using techniques such as batch normalization and weight regularization.

43. What is the concept of amortized inference in VAEs?

Answer: Amortized inference refers to the idea of learning a neural network to approximate the posterior distribution in a VAE.

44. What is the concept of variational inference in VAEs?

Answer: Variational inference is a method of approximating the posterior distribution in a VAE using a variational distribution.

45. How do you implement a VAE for text generation?

Answer: Text generation can be implemented using a VAE, where the encoder takes a text input and generates a latent code, and the decoder generates a new text output from the latent code.

46. What is the concept of disentanglement in VAEs?

Answer: Disentanglement refers to the ability of a VAE to separate the underlying factors of variation in the data, such as object pose, shape, and color.

47. How do you perform style transfer using VAEs?

Answer: Style transfer can be performed using a VAE, where the encoder takes a content image and a style image as input, and generates a latent code that combines the content and style.

48. What is the concept of attention in VAEs?

Answer: Attention refers to the ability of a VAE to focus on specific parts of the input data when generating output.

49. How do you implement a VAE for image compression?

Answer: Image compression can be implemented using a VAE, where the encoder takes an image input and generates a compressed latent code, and the decoder generates a reconstructed image from the latent code.

50. What is the concept of hierarchical VAEs?

Answer: Hierarchical VAEs are VAEs that use a hierarchical structure to model complex data distributions.

51. How do you implement a VAE for time series forecasting?

Answer: Time series forecasting can be implemented using a VAE, where the encoder takes a time series input and generates a latent code, and the decoder generates a forecasted time series output from the latent code.

52. What is the concept of normalizing flows in VAEs?

Answer: Normalizing flows are a type of invertible neural network that can be used to model complex data distributions in VAEs.

53. How do you implement a VAE for recommender systems?

Answer: Recommender systems can be implemented using a VAE, where the encoder takes a user-item interaction matrix as input and generates a latent code, and the decoder generates a recommended item output from the latent code.

54. What is the concept of amortized inference in GANs?

Answer: Amortized inference refers to the idea of learning a neural network to approximate the posterior distribution in a GAN.

55. How do you implement a GAN for data imputation?

Answer: Data imputation can be implemented using a GAN, where the generator takes a incomplete data input and generates a completed data

output, and the discriminator evaluates the generated data.

56. What is the concept of conditional GANs?

Answer: Conditional GANs are GANs that generate samples based on a specific condition or class.

57. How do you implement a GAN for video generation?

Answer: Video generation can be implemented using a GAN, where the generator takes a sequence of images as input and generates a video, and the discriminator evaluates the generated video.

58. What is the concept of attention in GANs?

Answer: Attention refers to the ability of a GAN to focus on specific parts of the input data when generating output.

59. How do you perform style transfer using GANs?

Answer: Style transfer can be performed using a GAN, where the generator takes a content image and a style image as input, and generates an output image that combines the content and style.

60. What is the concept of disentanglement in GANs?

Answer: Disentanglement refers to the ability of a GAN to separate the underlying factors of variation in the data, such as object pose, shape, and color.

61. How do you implement a GAN for text-to-image synthesis?

Answer: Text-to-image synthesis can be implemented using a GAN, where the generator takes a text input and generates an image, and the discriminator evaluates the generated image.

62. How do you implement a VAE for data augmentation?

Answer: Data augmentation can be implemented using a VAE, where the encoder takes an input image and generates a latent code, and the decoder generates a new image with augmented features from the latent code.

63. What is the difference between a confusion matrix and a classification report?

Answer: A confusion matrix is a table that summarizes the predictions against the actual true labels, while a classification report provides a summary of the precision, recall, and F1 score for each class.

64. How do you perform feature selection in machine learning?

Answer: Feature selection can be performed using techniques such as filter methods, wrapper methods, and embedded methods.

65. What is the concept of regularization in machine learning?

Answer: Regularization is a technique used to prevent overfitting by adding a penalty term to the loss function.

66. What is the concept of clustering in machine learning?

Answer: Clustering is a type of unsupervised learning where similar data points are grouped together based on their features.

67. What is the concept of a Generative Adversarial Network (GAN)?

Answer: A GAN is a type of deep learning model that consists of a generator and a discriminator, which are trained together to generate new samples that are similar to a given dataset.

68. What is the concept of a Variational Autoencoder (VAE)?

Answer: A VAE is a type of deep learning model that consists of an encoder and a decoder, which are trained together to learn a compressed representation of the input data.

69. How do you train a VAE?

Answer: A VAE can be trained using a probabilistic objective, such as the evidence lower bound (ELBO).

70. How do you implement a VAE for text-to-image synthesis?

Answer: Text-to-image synthesis can be implemented using a VAE, where the encoder takes a text input and generates a latent code, and the decoder generates an image from the latent code.

Written by Sujatha Mudadla

1.7K Followers

M.Tech(Computer Science),B.Tech (Computer Science) I scored GATE in Computer Science with 96 percentile.Mobile Developer and Data Scientist.

More from Sujatha Mudadla

Sujatha Mudadla

Interview questions on Android Jetpack Compose.

What is Android Jetpack Compose, and why was it introduced?

Jun 8, 2023 428 2

Sujatha Mudadla

Android Kotlin Coroutines interview questions

1. Q: What are Kotlin Coroutines, and how are they different from traditional threading?

Dec 19, 2023 701 2

Sujatha Mudadla

AWS Data Engineering Interview Questions.

Question: What is Amazon S3, and how is it commonly used in data engineering?

Jan 11 248 2

Sujatha Mudadla

Mobile Architect Interview Questions and Answers.

Question: What is the role of a Mobile Architect?

Dec 22, 2023 287 1

See all from Sujatha Mudadla

Recommended from Medium

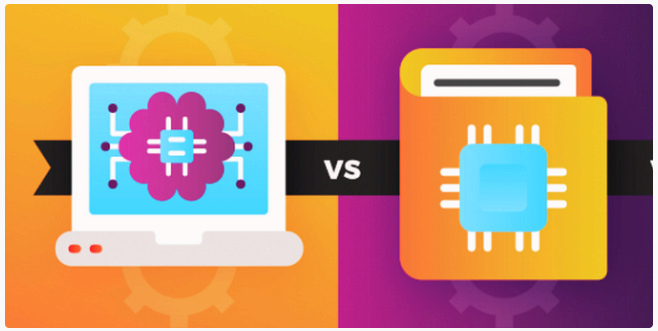

Michael Iyam

Punyakeerthi BL

AI Engineer vs. ML Engineer: Key Differences

The demand for artificial intelligence (AI) and machine learning (ML) engineers is high...

May 27 14 1

Difference between Self-Attention and Multi-head Self-Attention

Self-attention and multi-head self-attention are both mechanisms used in deep learning...

Apr 24 80

Lists

Staff Picks

742 stories · 1337 saves

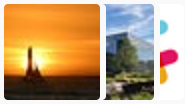

Stories to Help You Level-Up at Work

19 stories · 816 saves

Self-Improvement 101

20 stories · 2813 saves

Productivity 101

20 stories · 2399 saves

Mohammad Zeynali

Mastering Generative AI Interviews: Key Questions and Ho...

Introduction

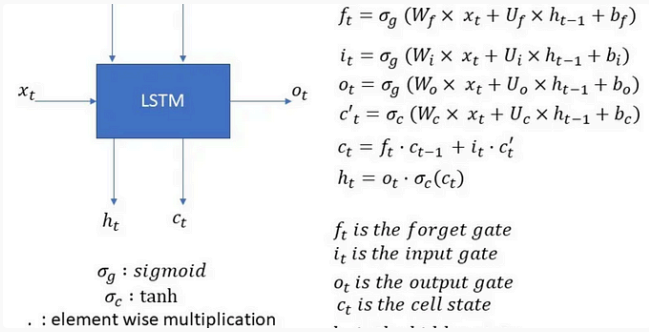

Milana Shkhanukova

100 NLP questions & ANSWERS | RNN| LSTM| GRU part

Today we will cover the questions regarding the RNN and LSTM.

Jun 10 27

Sep 11 38

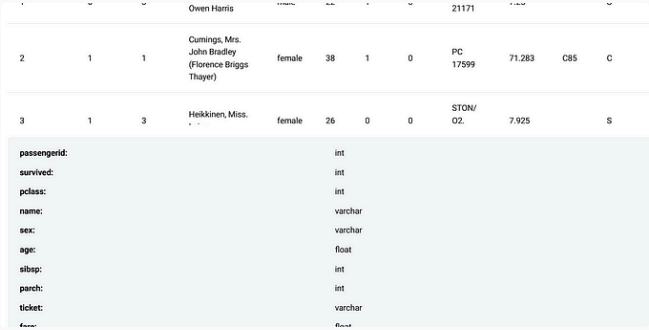

			Owen Harris					21171		
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38	1	0	PC 17599	71.283	C85	C
3	1	3	Heikkinen, Miss.	female	26	0	0	STON/ O2.	7.925		S

passengerid:	int
survived:	int
pclass:	int
name:	varchar
sex:	varchar
age:	float
sibsp:	int
parch:	int
ticket:	varchar

Harshit Singh in ILLUMINATION

SQL interview at Google

This question is published at stratascratch.com. Find the oldest survivor...

Sep 17 83 1

Sanjay Kumar PhD

Apache Spark Interview Questions and Answers

What is Apache Spark?

Sep 11 11

See more recommendations